

Par4All primer with tpips

Amaury DARSCH

HPC Project

October 11, 2010

Contents

- 1 General concepts
- 2 First contact
- 3 Working with OpenMP
- 4 Special operations
- 5 Design flow

Overview

- Par4All is a framework designed for automatic program parallelization
- Heavily relies on PIPS source-to-source compiler framework
 - Programming Integrated Parallel System
 - In French: *Parallélisation Interprocédurale de Programmes Scientifiques*
- Par4All provide a p4a script for easy parallelization...
- ...But some people may want to use more expressive interfaces: tpips, PyPS, ipyps...
- This document describes basic tpips concepts and usage

Installation and documentation

- Installation methods

<http://www.par4all.org/download/RELEASE-NOTES.txt>

- Documentation

<http://www.par4all.org/documentation>

- Compilation by your own

- `git clone git://git.hpc-project.com/git/par4all.git`

- `cd par4all`

- `git checkout -b p4a remotes/origin/p4a`

- `./src/simple_tools/p4a_setup [--prefix]`

- Activation

`source run/etc/par4all-rc.sh`

Definitions of PIPS concepts

- A module is the minimum processing unit of a program to work on
 - A function or subroutine
 - File global definition are in a special module: a compilation unit
 - A source file is split in a collection of modules
- A workspace is a collection of modules from several files
 - A workspace is created with the *create* command
 - Modules are mainly created when files are added to a workspace
- A user transformation is often a succession of PIPS phases
 - A phase is an elementary operation on resources (code, dependence graph, symbol tables, call graph...)
 - A resource is produced by a phase
 - The initial resources are automatically created

Module naming conventions

Source file

ex.c

```
// function f1
int f1 () {
    return 1;
}
```

```
// function f2
static int f2 () {
    return 2;
}
```

Module f1

f1

```
int f1 () {
    return 1;
}
```


Module f2

ex!f2

```
static int f2 () {
    return 2;
}
```

Basic flow

- Split & parse
Split the source files in modules and generate internal representation
- Analyze & transform
Perform one or several code analysis and transformations
- Unsplit
Regenerate the target source code from the database



Testing your installation

- Set the Par4All path and environment variable
 - `bash$ source run/etc/par4all-rc.sh`
- Start the tpips command interpreter
 - `bash$ tpips`

Session example

```
bash$ tpips
tpips (ARCH=SOFT_ARCH)
running as tpips
...
tpips>quit
bash$
```


First example

C function

ex1.c

```
long max (const long x, const long y) {  
    return x < y ? y : x;  
}
```

Command file

ex1.tpips

```
# Just in case there was already an ex1 workspace:  
delete          ex1  
create          ex1 ex1.c  
display         PRINTED_FILE(max)  
quit
```

First example

Session

Shell session

ex1.tpips

```
bash> tpips ex1.tpips
tpips (ARCH=LINUX_x86_64_LL)
  running as tpips
...
PRINTED_FILE made for max.
long max(const long x, const long y)
{
    return x<y?y:x;
}
```

First example

Commands

Command file		<i>ex1.tpips</i>
delete	ex1	
create	ex1 ex1.c	
display	PRINTED_FILE(max)	

- *delete ex1*
Remove the *ex1* workspace
- *create ex1 ex1.c*
Create a new *ex1* workspace and adds the source file *ex1.c*
- *display PRINTED_FILE(max)*
Display the module *max* stored in the database

First example

Database

Command file		<i>ex2.tpips</i>
delete	ex1	
create	ex1 ex1.c	
apply	UNSPLIT(%ALL)	

- *create ex1*

Create a workspace as a *ex1.database* directory

- *apply UNSPLIT(%ALL)*

Rebuild the source files after one or several transformations

The files are placed in the *ex1.database/Src* directory

Generic commands

- *create name files...*
Create a new workspace by adding source files
- *delete name*
Delete a workspace by name
- *close*
Close a workspace
- *echo message*
Print a message in the terminal window
- *shell command*
Execute a shell command
- *quit*
Exit the interpreter (*exit* also valid)
- *# some comments...*

Resources and phases commands

- *module name*

Select a module by name and make it the default one

- *apply phase*

Execute a phase (transformation...) on the current module

- *apply phase(name)*

Execute a phase on the specific named module

- *apply phase(%ALLFUNC)*

Execute a phase on workspace functions

- *activate rule*

Activate a particular rule when there are many ones to build a resource

- *display resource*

Display a resource by name.

Module references

- *%ALLFUNC*
All functions or subroutines in the workspace
- *%ALL*
All modules in the workspace, that are *%ALLFUNC* plus all the compilation units describing source files
- *%MODULE*
The current module
- *%CALLEES*
All modules called in the given module
- *%CALLERS*
All modules that calls the given module
- *%PROGRAM*
The current program

Par4All and OpenMP

- Support for OpenMP 2.5
 - Automatic pragma insertion in the code
 - Need to compile with the `-fopenmp` option with GCC
- Fine Grain Parallelization
 - Based on ALLEN & KENNEDY algorithm
 - Mostly for inner loops
 - Can find more parallelism by distributing everything
 - Put the pressure on memory interface
- Coarse Grain Parallelization
 - Based on array regions analysis
 - Mostly for outer loops
 - No loop distribution
- Both methodologies are supported
 - Difficult to choose between one or the other
 - Experiments are essential in understanding performance gains

Loop parallelization

Initial code

```
void tinit (const long size, long t[size]) {  
    long i = 0;  
    for (i = 0; i < size; i++) t[i] = 0L;  
}
```

Loop parallelization

```
#pragma omp parallel for  
    for(i = 0; i <= size-1; i += 1)  
        t[i] = 0L;
```

OpenMP Generation

PIPS directives

```
# Privatize loop local scalar variables:
```

```
apply      PRIVATIZE_MODULE(tinit)
```

```
# Apply parallel code transformation
```

```
apply      INTERNALIZE_PARALLEL_CODE(tinit)
```

```
# Generate OpenMP directives
```

```
apply      OMPIFY_CODE(tinit)
```

Fine grain parallelization — complete example

PIPS script

```
delete      omp
create      omp omp.c
# Select plain style source output:
setproperty PRETTYPRINT_SEQUENTIAL_STYLE "do"

apply      PRIVATIZE_MODULE(tinit)
apply      INTERNALIZE_PARALLEL_CODE(tinit)
apply      OMPIFY_CODE(tinit)
# Regenerate in .../Src all the transformed files:
apply      UNSPLIT(%ALL)
close
quit
```

Coarse grain parallelization — complete example

PIPS script

```
delete      omp
create      omp omp.c
setproperty PRETTYPRINT_SEQUENTIAL_STYLE "do"

apply       PRIVATIZE_MODULE(tinit)
apply       COARSE_GRAIN_PARALLELIZATION(tinit)
apply       OMPIFY_CODE(tinit)
apply       UNSPLIT(%ALL)

close
quit
```

Coarse grain array swapping

Array Swapping

```
void tswap (const long size,  
            double x[size], double y[size]) {  
    long    i = 0;  
    #pragma omp parallel for  
    for (long i = 0; i < size; i++) {  
        double t = x[i];  
        x[i] = y[i];  
        y[i] = t;  
    }  
}
```

Private variable example

- The variable `t` has been privatized by PIPS because declared outside of the loop

Array Swapping

```
void tswap (const long size,  
            double x[size], double y[size]) {  
    long i = 0;  
    double t;  
    #pragma omp parallel for private(t)  
    for (long i = 0; i < size; i++) {  
        t = x[i];  
        x[i] = y[i];  
        y[i] = t;  
    }  
}
```

Call graph

PIPS Script

```
delete      cg
create      cg graph.c

display     CALLGRAPH_FILE(%ALL)

close
quit
```

Dependence graph

PIPS Script

```
delete      dg
create      dg graph.c

# Compute more precise dependences by using array regions
activate    REGION_CHAINS
# Ask for exact regions when possible
activate    MUST_REGIONS

display     DG_FILE(%ALL)

close
quit
```


Automated flow

- Project structure:
 - All transformations integrated in a project build
 - Based on *Makefile*
 - Add a rule to build with PIPS optimization
- File access:
 - All files located in the database (*Src* directory)
 - Compilation in place
 - Must be cautious with dependencies
- Project testing:
 - Must have regression
 - Must compare with sequential results
 - Must exhibit performances benefits

Going further

- There are around 300 phases in PIPS ! ☺
- Look at the documentation (such as *PIPS analyses and transformations*) on
<http://pips4u.org/doc/pips-technical-pages> and
<http://par4all.org>
- Automate stuff by using PyPS Python front-end
- p4a can accept some Python injection
http://download.par4all.org/doc/simple_tools/p4a
- Ask HPC Project for consultancy and training ☺